

1 Tree Based Models

The idea behind this family of methods is to segment the region into subspaces and the prediction for any subspace is made by taking the *mean* or *mode* of the response in that region.

A simple decision tree may not give performance at par with linear regression or generalized additive models, but when combined with techniques like bagging and boosting, multiple trees can yield significant reduction in bias.

A simple decision tree finds its utility in being relatively simpler in interpretation than its derived non-linear family.

1.1 Regression Trees

The algorithm behind building trees consists of two basic steps

1. Divide the predictors X_1, X_2, \dots, X_n into J **distinct** and **non-overlapping** regions R_1, R_2, \dots, R_J .
2. For making a prediction in the region R_j , simply take the **response mean** of all the data points following in that region.

With the above definition, the loss becomes

$$RSS = \sum_{j=1}^J \sum_{i \in R_j} (y_i - \bar{y}_{R_j})^2$$

But minimizing this loss is infeasible in practice as the total number of possible partitions are too large !

To overcome this, we build the trees in a greedy top down approach called **recursive binary splitting**. This method aims to find the best possible split at each level in a greedy manner.

1.1.1 The Algorithm

To start off, the algorithm will try to split the data into two regions $\{X|X_j < s\}$ and $\{X|X_j \geq s\}$ using the values of predictor X_j and the search is performed across all the predictors to choose the one which minimizes the RSS. Mathematically

$$R_1(j, s) = \{X|X_j < s\} \text{ and } R_2(j, s) = \{X|X_j \geq s\}$$
$$\underset{j,s}{\text{minimize}} \sum_{i:x_i \in R_1} (y_i - \bar{y}_{R_1})^2 + \sum_{i:x_i \in R_2} (y_i - \bar{y}_{R_2})^2$$

where \bar{y}_{R_1} and \bar{y}_{R_2} are the response mean in the region R_1 and R_2 respectively.

The same algorithm is run recursively. The only change that is made that after the first split, we now work independently on two separate regions. For each region, we will not be able to consider all the possible values of s since the region has been restricted. Otherwise, the RSS term will remain the same.

The process is stopped when a pre-determined stopping criteria is reached. For making a prediction, we will first determine the region in which the test observation falls and then make the prediction based on the mean of the training samples in this region.

1.1.2 Tree Pruning

The tree build using above strategy can have good results on the training set, but will yield poor performance on the test set due to the high complexity. A better strategy is to build **short trees which are less complex and produce lower variance at cost of little bias** and are more interpretable.

A simple strategy can be to build only when there is a decrease in RSS above a threshold. This approach can be short-sighted as some future good branches can be missed.

Tree pruning is to **build a large tree and then shrink it back to obtain a subtree**. The RSS of a subtree can be found via cross validation. This approach is expensive for all possible subtrees.

Cost complexity Pruning or **weakest link pruning** is an approach to overcome this problem. For a non-negative parameter α , there exists a subtree T of the large tree T_0 such that the following is minimized

$$RSS = \sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \bar{y}_{R_m})^2 + \alpha |T|$$

where $|T|$ is the number of terminal nodes in the tree. The formulation is similar to lasso and here we are controlling the complexity of the tree via a parameter α . Larger the α , the less number of terminal nodes there would be in the tree.

The algorithm can then be summarized as follows

1. Use *binary recursive splitting* to obtain a large tree on the data set. Stop only when at a terminal node, the number of observations is less than a minimum.
2. Apply *cost complexity pruning* in order to obtain a small tree T as a function of the cost parameter α (fix a set of α and get a set of subtrees).
3. Use K-fold cross validation to choose α . For each of the folds $k = 1, 2, \dots, K$
 - (a) Repeat steps 1 and 2 excluding the data from the k th fold to obtain a family of subtrees as a function of α .
 - (b) Evaluate the mean predicted validation error as a function of α .
Choose the value of α that minimizes the average error across all the folds.
4. Select the subtree from step 2 that corresponds to the chosen value of α .

1.2 Classification Trees

Here we predict a qualitative variable and will report the **most occurring class in the region**. Since we are almost always also interested in the probabilistic aspects of the prediction, we also consider **the proportion of the classes occurring in a region**.

1.2.1 The Algorithm

The algorithm is almost similar to the regression trees but instead of RSS, we need to use something that is related to class proportions and frequencies. A very straightforward metric is **classification error rate**

$$E = 1 - \max_k \hat{p}_{mk}$$

where \hat{p}_{mk} is the proportion of the k th class in the m th region. However, in practice, this metric is not sufficient to grow the trees.

1.2.2 Gini Index

Gini Index is defined as

$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

for a given region. Clearly, G is close to zero when the node is pure, i.e. most of the observations come from the same class. Gini index is also known as a measure of **purity** since its low value implies predominance of a single class in the node.

1.2.3 Entropy

Entropy is defined as

$$D = \sum_{k=1}^K -\hat{p}_{mk} \log(\hat{p}_{mk})$$

for a region. Entropy takes value close to zero if the class probabilities are close to 0 or 1. Gini index and Entropy are somewhat similar numerically.

1.3 Bagging

Decision trees suffer from a high variance. When the data is divided into smaller parts, it is quite likely to expect variance when the data set itself is changed. Bootstrap Aggregation or Bagging is a method to try to overcome this problem.

Averaging a set of random variables reduces variance. Consider n random variables with the same variance σ^2 , then the variance of their average is σ^2/n . Naturally, we can extend this idea to build several prediction models on different training data sets and average out their results to reduce the variance in the response.

We train b different models, make predictions $\hat{f}_i(x)$ using these and then average them out as

$$\hat{f} = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x)$$

to get a low variance prediction statistic.

Since we do not have access to different training datasets, we will get the average using bootstraps of the original data set. Note that **trees grown on bootstrapped data sets are deep and not pruned** so that they have low bias. Even though they may have high variance, averaging will reduce it out.

Averaging bootstrapped predictors works for regression while we take the majority vote in case of classification. The overall prediction is the most commonly occurring class across the B trees. Note that **using a large value of B will not lead to overfitting**, it should just be sufficiently large to ensure the error has come down compared to a single tree (building too many trees can take up significant time).

1.3.1 Out of Bag Error

It can be shown the probability of an observation being present in a bootstrapped dataset is $1 - 1/e$ (this can be shown by considering that the probability that an observation is not present in the data set is $((1 - 1/n)^n$ and taking the limit $n \rightarrow \infty$). Thus, on an average, a data will only contain about 2/3rd (0.63 to be exact) of the total observations in it. We can make use of the observations not used for training to estimate the error.

Any observation not part of the bootstrapped set is called **Out of Bag** and the error we are about to calculate is out of bag error. Now, we can take the i th observation, and get its prediction from the sets where it was not used for training. This will be around 1/3rd of the predictors. We can combine these results using average in case of regression or majority vote in case of classification. Thus, we have a "test error" for all the observations in the data.

1.3.2 Variable Importance

We have been able to reduce the variance using the bagging approach, but the model is no longer interpretable due to the presence of multiple trees. We need to somehow aggregate all the trees to get this measure.

A simple workaround is to calculate the total amount the RSS (regression) or Gini Index (classification) has decreased across all the trees due to split on a particular variable. To make the values comparable, we simply take the average across B trees. A large value means that the variable has high importance.

1.4 Random Forests

Random forests overcome the problem of **decorrelation** in bagging. The trees built using bagging are similar because if there is a strong predictor, we expect it to dominate the splits in each of the trees and give similar trees.

Note that adding correlated variables will only change the variance of the mean slightly. Consider $Var(X + Y)$ which is $Var(X) + Var(Y)$ when X and Y are independent and $Var(X) + Var(Y) + 2Cov(X, Y)$ when X and Y are not independent. Clearly, the variance of the mean will be close to individual variances when the variables are not independent. Hence, to improve the classifier, the predictions must not be correlated or, the trees themselves must be decorrelated.

Variance of the mean of independent variables is much less than their individual variances while we don't expect such significant reduction in the case of correlated variables.

Random Forests achieve decorrelation by using only a subset of variables for each split (\sqrt{p}). This subset is chosen randomly at each split to ensure that the trees differ enough in structure to give less correlated predictions.

Similar to bagging, **random forests will not overfit if we increase the number of trees B .** Hence, the choice can be made based on whether the reduction in test error is sufficient.

1.5 Boosting

Boosting is a general technique that can be applied to multiple statistical learning techniques and not just decision trees. Like Bagging, boosting is also trained on multiple data sets derived from

the training data, but instead of building independent trees on different data sets, **boosting is a sequential process that builds trees on modified versions of the original data set.**

Given a model, boosting will try to fit a tree on the residuals obtained by the model instead of the response Y . The model is then updated by adding this tree and the residuals are calculated again. **Trees in boosting are typically small, just a few terminal nodes** controlled by the parameter d and we allow the tree to fit the areas where improvement is needed in a slow manner. Shrinkage parameter λ allows to further slow down this process allowing for a variety of trees.

Boosting has three parameters

1. B , the number of trees. Unlike bagging, large B can overfit although slowly. Right value of B can be determined through cross validation.
2. Shrinkage parameter λ which controls how slowly the model is learnt and is a small positive number in the range 0.1 to 0.001 typically. Very small λ can require a very large B to achieve a good performance.
3. The number of splits d in a single tree. Typically $d = 1$ is used and referred to a stump. Higher value signifies higher interaction between the variables. When using stumps, the model is simply additive since we are using a single variable in each of the trees.

Boosting algorithm is as follows

1. Set the residuals $r_i = y_i$ and the model $\hat{f}(x) = 0$.
2. For $b = 1, \dots, B$ repeat
 - (a) Fit the model $\hat{f}^b(x)$ with d splits on the data (X, r)
 - (b) Update the model

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$$

- (c) Update the residuals

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x)$$

3. Return the boosted model

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x)$$